

[« Prev](#)[Next »](#)

Database Weekly

Issue 68 — August 21, 2015

FEATURED

How Does A Relational Database Work — It sounds basic, but this article elegantly dives into the performance of various data structures behind databases, shows how relational databases perform under the hood, and even goes into query parsers, optimizers, and locking strategies.

CODING GEEK

ORMs Under The Hood — A look at what ORM (Object Relational Mappers) are, what they offer, their relationship to SQL, and a summary of the arguments for and against using them.

PATRYCJADYBKA

Azure DocumentDB: JavaScript as Modern Day T-SQL — Andrew Liu reflects on the DocumentDB team's vision for JavaScript developers to work natively using JS and JSON without entity mappers, schemas, and similar 'duct tapes'.

MICROSOFT AZURE BLOG

Cloud-Scale Database Monitoring Made Easy with Datadog — Use Datadog's turn-key integrations to seamlessly aggregate metrics and events across your databases and entire devops stack. Gain powerful insight into all your apps, tools and services. Start your Free Datadog Trial today!

DATADOG **SPONSORED**



CurioDB: A Distributed, Persistent Redis Clone — CurioDB is a distributed and persistent Redis clone, built with Scala and Akka. It's currently considered a 'toy project' by its creator, however.

STEPHEN McDONALD **CODE**

'Big Data' Features Coming in PostgreSQL 9.5 — Postgres 9.5 will be here soon and here's a rundown of the features that help with dealing with big data and performance.

LWN

Comparing MongoDB Performance On Public Clouds: AWS, Azure & Digital Ocean — Digital Ocean isn't really comparable to AWS and Azure IMHO, but the performance is great.

MONGODIRECTOR

Keyvi: A New Highly Efficient, Scalable Key Value Index — keyvi is an immutable, FST-based key value store with high compression ratio and extreme scalability.

CLIQZ

The SQLite Query Optimizer Overview — “This document provides overview of how the query planner and optimizer for SQLite works.”

SQLITE

Labellio: Scalable Cloud Architecture for Efficient Multi-GPU Deep Learning — A look at a cloud-based GPU oriented deep learning system focused on computer vision applications.

PARALLEL FORALL

Extending Seven Bridges Genomics with Amazon Redshift and R — Redshift is AWS’s massive parallel processing data warehouse service. This article takes a end-to-end walkthrough of performing genomics analysis on the system.

AMAZON

IN BRIEF

[AWS Data Pipeline Support for all Amazon RDS Databases Now Available](#)

AMAZON **NEWS**

[MongoDB Announces Support for IBM z Systems’ Mainframe Solutions](#)

MONGODB **NEWS**

[Long Live Your App: MongoDB Replication](#)

MARCO GALLARDO

[Become A MySQL DBA Blog Series: The Query Tuning Process](#)

SEVERALNINES

[Using SQLite4's LSM Storage Engine as a Stand-alone NoSQL Database with Python](#)

CHARLES LEIFER

[The 8 Best Ways To Lose Your DBA](#) — To stop churning through DBAs, don’t make your DBA’s a single point of failure and avoid bureaucracy, amongst other things.

VMDCORTEX

[SqlSmash V2: Faster and Better](#) — SqlSmash is a plugin for SQL Server Management Studio that helps you write maintainable SQL scripts and work faster.

DOTNETSURFERS **TOOLS**

[The Amazon DynamoDB Storage Backend for Titan](#)

AMAZON WEB SERVICES **CODE**

[mongo-aggregation-debugger: An Aggregation Pipeline Debugger for MongoDB](#) — Debug MongoDB’s aggregation framework and visualize what each stage of the pipeline outputs.

GUILLAUME FLANDRE **CODE**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).