

[« Prev](#)[Next »](#)

## Database Weekly

Issue 72 — September 25, 2015

### FEATURED

**Scylla DB: An Apache Cassandra-compatible Database Written in C++** — From the team behind KVM comes a new open source database that promises to be the ‘world’s fastest column-store database’. It’s Apache Cassandra compatible and can just be dropped in.

SCYLLA

**SQL vs NoSQL: How to Choose** — A look at some different scenarios and applying the right criteria to choose whether the SQL or NoSQL route is right for them.

CRAIG BUCKLER

**9 Tips for Better Database Design** — Plan ahead, document your model, follow conventions, think carefully about your keys, use integrity checks carefully, and more.

JEFFREY EDISON

**Real-time Database Monitoring at Any Scale with Datadog** — Use Datadog's turn-key integrations to seamlessly aggregate metrics and events across your databases and entire devops stack. Gain powerful insight into all your apps, tools and services. Begin your Free Datadog Trial today!

DATADOG **SPONSORED**



**Handling Circular Data Structures in Postgres** — If you're dealing with groups that contain other groups, or any other kind of nested structure the default might be to simply make a bunch of separate database calls from within your app. Fortunately with recursive queries there is a better way.

GARY SIELING

**A New Breed of Database Hopes to Blend The Best of NoSQL and RDBMS** — Multi-model databases may help tame the growing complexity of enterprise data. Matt Asay explains.

TECHREPUBLIC

**Yesquel: Scalable SQL Storage for Webapps [PDF]** — “We design, develop, and evaluate Yesquel, a system that provides performance and scalability comparable to NoSQL with all the features of a SQL relational system.” An academic paper.

AGUILERA, LENERS, WALFISH

### IN BRIEF

[Elasticsearch 2.0.0-beta2 Released](#) — The final Elasticsearch 2.0 is just around the corner.

ELASTIC **NEWS**

[Enterprise Cassandra Gets a Storage Boost](#)

ENTERPRISE TECH **NEWS**

[Inside Platfora's Transition to Apache Spark](#)

DATANAM **NEWS**

[Tape Storage Spec Jumps to 15TB](#) — You've got to keep that data backed up somewhere.

DATANAM **NEWS**

[Better Density and Lower Prices for Azure's SQL Elastic Database Pools](#)

SCOTT GUTHRIE **NEWS**

[Working with Email Addresses in SQL](#) — Check if an email is valid (with a regexp - be careful!), split an email into its key parts, and more, all from SQL.

BENN STANCIL

[Real-Time Search With MongoDB and Solr](#)

YANG ZHAO

[Delightful Database Seeding with Docker](#) — “Using Docker we have created a user friendly tool that anyone in our organization can use, (even sales) both when resetting a demo laptop, and when getting a new on-site installation quickly up and running.”

KRISTIAN HELGESEN

[Using the SQLite JSON Extension with Python](#)

CHARLES LEIFER

[Storing Data with Redis](#) — “Would you stuff all of your data into one database table? That’s crazy, Mike, don’t be silly! What if I told you most people do just that with Redis?”

MIKE PERHAM

[The Simple Power of Elasticsearch Analyzers](#)

DOUG TURNBULL

[The SELECT Problem in Redis](#) — “The need to maintain your own indexes for structured data is the sole reason that makes Redis unsuitable for certain applications.”

FEDERICO PEREIRO

[How Medium Detects Hotspots in DynamoDB using Elasticsearch, Logstash and Kibana](#)

MEDIUMENGINEERING

[Deploy An Elastic High-Availability SQL Cluster with Crate and Weave](#)

GIANT SWARM **TOOLS**

[PolyBase in SQL Server 2016](#) — PolyBase is a new feature that integrates SQL Server Parallel Data Warehouse (PDW) with Hadoop.

CHANNEL 9 **TOOLS** **VIDEO**

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).