

[« Prev](#)[Next »](#)

Database Weekly

Issue 73 — October 2, 2015

FEATURED

Amazon Launches 'Elasticsearch Service' (ES) — Elasticsearch is a real-time, distributed search and analytics engine (which effectively acts like a document database too) and Amazon have launched a new service around it.

AWS OFFICIAL BLOG

Rethinking NoSQL: 4.5 Reasons Why RethinkDB Is Worth a Look — RethinkDB is redefining NoSQL, in a good way. In this talk Rob Conery will show you the most compelling features of RethinkDB 2.0, and why it deserves consideration for your next project.

ROB CONERY 

SQLite As An Application File Format — In many cases, SQLite is a better choice than either a custom file format, a pile-of-files, or a wrapped pile-of-files. SQLite is a high-level, stable, reliable, cross-platform, widely-deployed, extensible, performant, accessible, concurrent file format.

SQLITE

Get more from Redis: Autoscaling and Performance Dashboard from RedisGreen — Get insights quickly with our analytics dashboard. Track usage and quickly identify performance bottlenecks. All plans backed by support from engineers who have been scaling Redis for years.

REDISGREEN 



Introducing U-SQL: A Language For Easy Big Data Processing — A language that's part of Azure Data Lake Analytics that unifies the benefits of SQL with the expressive power of your own code. U-SQL's scalable distributed query capability enables you to analyze data across relational stores such as Azure SQL Database.

MICROSOFT

Clarifications About Redis and Memcached — In response to [Storing Data with Redis](#) where Mike Perham recommended memcached over Redis for caching, Redis's creator clarifies why Redis really is a good choice for caching.

SALVATORE SANFILIPPO

Lazy Redis is Better Redis — The creator of Redis shares the complex tale of implementing a non-blocking DEL (delete) command for data stored in a Redis instance.

SALVATORE SANFILIPPO

NoSQL LinkedIn Skills Index for September 2015 — “Once again MongoDB was responsible for more than 50% of all mentions of NoSQL database in LinkedIn member profiles in Q3, placing it way, way ahead of the nearest competitor.”

MATTHEW ASLETT

Modern SQL: The Inaugural Post — Markus Winand, the author of *Use The Index, Luke*, has launched Modern-SQL.com which has the aim of helping you explore modern SQL (SQL:2011 for now).

MARKUS WINAND

The Way to Open Source Hadoop Native SQL — Pivotal has open sourced [HAWQ](#) and [MADlib](#), contributing them to the Apache Software Foundation meaning there is now a full featured, SQL standards compliant, interactive SQL engine for Hadoop available in open source.

PIVOTAL [TOOLS](#)

JOBS

[Golang Database/Distributed Systems Engineer at InfluxDB \(anywhere\)](#) — If you've got a background building databases or distributed systems and want to store and process data at scale using Go, please apply.

INFLUXDB

IN BRIEF

[SQL Server 2016 Community Technology Preview 2.4 Is Available](#)

MICROSOFT [NEWS](#)

[Google Launches Cloud Dataproc, A Managed Spark And Hadoop Big Data Service](#) — The new Google Cloud Dataproc service sits between the Spark data processing engine or Hadoop framework directly on VMs and a fully managed service like Cloud Dataflow.

TECHCRUNCH [NEWS](#)

[How Will New Memory Technologies Impact In-Memory Databases?](#)

HIGH SCALABILITY

[Architecture of SQLite](#)

SQLITE

[RDBMS Genealogy: A PDF of the History of Databases](#)

HASSO-PLATTNER-INSTITUT

[Petabyte-Scale Data Pipelines with Docker, Luigi and Elastic Spot Instances](#)

ADROLL

[Scaling Writes on Amazon DynamoDB Tables with Global Secondary Indexes](#)

AWS BIG DATA BLOG

[SQL vs NoSQL: The Differences](#)

SITEPOINT

[FreeGeoDB: Free Database of Geographic Place Names and Corresponding Geospatial Data](#)

DELIGHT.IM [CODE](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).