

[« Prev](#)[Next »](#)

## Database Weekly

Issue 87 — January 15, 2016

### FEATURED

**Designing Schemaless, Uber Engineering's Scalable Datastore** — Uber has developed its own MySQL-based schemaless storage system. This article explains why and how it came together.

UBER ENGINEERING

**Writing More Legible SQL** — Writing SQL is one thing, but when it comes to reading other's SQL, things can get tricky. These few tips will help to make your SQL much more understandable by others.

CRAIG KERSTIENS

**Real-time Postgres monitoring** — Visualize Postgres system resources, connections, locks, transactions, and indexes, and instantly correlate with metrics from anywhere in your stack.

DATADOG **SPONSORED**



**Using CTEs and Unions to Compute Running Totals** — A look at combining CTEs with unions to give not just grouped results, but also side by side aggregates.

PERISCOPE DATA.COM

**How Developers Use Databases Today: A JetBrains Report** — Insights based on responses from 2000 developers. One factoid: Postgres is the database most would like to migrate to, while MySQL is the one they most want to leave behind.

JETBRAINS

**Elasticsearch Replaces Solr As The Most Popular Search Engine** — While they're commonly used for their search functionality only, Solr and Elasticsearch are also document database engines of sorts. Elasticsearch has now pulled ahead of Solr in DB Engines' rankings.

DB ENGINES

**Querying Amazon Kinesis Streams Directly with SQL and Spark Streaming** — What if you could use your SQL knowledge to discover patterns directly from an incoming stream of data? Streaming analytics is a very popular topic of conversation around big data use cases.

AMAZON

**PipelineDB Releases First Enterprise Version Of Its Streaming SQL Database** — PipelineDB is built on a new way of looking at SQL databases, thinking about streams of data rather than data at rest in big silos.

TECHCRUNCH

**Hadoop's 10th Birthday** — A look back at how Hadoop has gone from an unknown to the king of distributed data set processing systems.

CLOUDERA

**Creating a Data Model For Carpooling** — It's always good to see how a database designer approaches modelling a situation. In this case, it's a ride-sharing app.

SHANTANU KHER

IN BRIEF

[Does your company need Big Data? LambdaNow can help!](#) — LambdaNow solves the complexity of Big Data by providing easy Big Data processing pipelines. Try it for FREE Now.

LAMBDA NOW **TOOLS** **SPONSORED**

[SQLite 3.10.1 Released](#) — ‘Fixed a 16-month-old bug in the query planner that could generate incorrect results’

SQLITE **NEWS**

[Picking the Right SQL-on-Hadoop Tool for the Job](#)

DATANAM **NEWS**

[Preview The Newest ODBC SQL Server Driver for Windows and Linux](#)

SQL SERVER TEAMBLOG **NEWS** **TOOLS**

[Survey: Big Data Goes Mainstream](#) — “Financial services firms remain the biggest users of big data technologies”

DATANAM **NEWS**

[JSON Support in SQL Server 2016](#)

CODINGSIGHT **NEWS**

[Top 6 Database Performance Metrics to Monitor in Enterprise Applications](#)

OMED HABIB

[What's New in Cassandra 3.0 - Part 1](#)

SMARTCAT

[Using the Ltree Datatype in Postgres](#)

LEAPFROG ONLINE

[Natural Language Processing with Spark on Reddit Comments](#)

DATAREFLEXIONS

[Graph Databases for Beginners: ACID vs. BASE Explained](#)

NEO4J

[As Data Science Evolves, It's Taking Statistics with It](#)

DATANAM

[PostgreSQL 9.6 Wish List](#)

NEIL MCGUIGAN

[Medis: An Open Sourced Redis GUI Tool for OS X](#) — Lets you view and edit keys, supports SSH tunnelling, config viewing/editing. It's built using Node, Electron and React.

ZIHUALI **TOOLS**

[MarsDB: A Promise Based Lightweight Database with MongoDB Query Syntax, Written On ES6](#)

ARTEMARTEMEV **CODE**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).