

[« Prev](#)[Next »](#)

Database Weekly

Issue 88 — January 22, 2016

FEATURED

dbbench: Bringing Active Benchmarking to Databases — dbbench is a fast, lightweight database workload generator that executes a workload defined a flexible configuration file.

ALEX REECE

Engine.so: A Catalog of Embeddable Data Storage Engines — A quick list of 21 different embeddable databases from obvious candidates like SQLite and Berkeley DB to unusual ones like Rlite and UpscaleDB.

DMTRY SIMONENKO [TOOLS](#)

OneCache: A Replicated Key/Value Store Accessible Via The Memcached Protocol — OneCache nodes natively support clustering and best effort replication. Adding nodes into the cluster effectively increases the in-memory storage capability of the cache.

ALEX DADGAR [CODE](#)

PostgreSQL Past, Present, and Future: Moving the Goalposts —

Learn about PostgreSQL: the history, the present, and the future from Robert Haas, Chief Architect, Database Server at EnterpriseDB and PostgreSQL Major Contributor and Committer.

ENTERPRISEDB [SPONSORED](#)



Business Logic in The Database. Yes or No? It Depends — Why is it so controversial to have business logic in the database? In which situations is it truly the right place and when is it not? Lukas Eder investigates.

LUKAS EDER

MySQL to Amazon Redshift Replication: A How-To Guide — Why do companies replicate their MySQL databases to Amazon Redshift and how can you approach it yourself?

ITAMAR WEISS

Impress Your Coworkers by Using SQL UNPIVOT — Some advanced SQL wrangling indeed.

JOJOQ

Adding Columns with Default Values to Really Large Postgres Tables from Rails

— For most frameworks it's common for you to set a default value when you add a new column, the problem is this can cause a full table re-write and hold a lock—read: not good. There's a few methods to work around this.

WEWORK ENGINEERING

IN BRIEF

[Unlimited Scalability & Unprecedented Speed with In-Memory Computing](#) — GridGain enables high-performance transactions, real-time streaming and fast analytics using a distributed, massively parallel architecture.

GRIDGAIN IN-MEMORY DATA FABRIC **TOOLS** **SPONSORED**

[Survey Sees Spark Emerging in 2016](#)

DATANAM **NEWS**

[MariaDB Raises \\$9M More, Michael Howard Named New CEO, Monty Widenius CTO](#) — The company behind the 'successor' fork (of sorts) of MySQL continues to march ahead.

TECHCRUNCH **NEWS**

[Hadoop Market is Neck and Neck, Forrester Says](#) — Forrester 'found four strong performers in a market it says is neck and neck.'

DATANAM **NEWS**

[One Approach to Boosting Data Accuracy at Scale](#)

DATANAM **NEWS**

[What Does An Oracle Audit Look Like? This One Certainly Wasn't Pretty](#) — Mars Inc. filed a lawsuit to get Oracle to scale back its license auditing efforts.

COMPUTERWORLD **NEWS**

[How to Install MongoDB 3 on Ubuntu 14.04](#)

MUHAMMAD PANJI

[Finding Nearest Neighbors in SQL](#) — A look at using Pythagoras' Theorem to calculate distances in SQL.

PERISCOPE

[Getting Started with JSON Support in SQL Server 2016](#)

DATABASE JOURNAL

[Testing Apache Cassandra with Jepsen](#)

DATASTAX

[SQL or NoSQL](#)

SHAULI GAL

[schemazen: Script and Create SQL Server Objects Quickly](#)

SETH RENO **TOOLS**

[Rethinking SIMD Vectorization for In-Memory Databases \[PDF\]](#) — An academic paper.

POLYCHRONIOU, RAGHAVAN AND ROSS

[Real-time Amazon Aurora performance monitoring](#) — Collect and aggregate metrics from across your Amazon Aurora cluster with ease. Correlate Aurora metrics with metrics from the rest of your stack. [Try it today with a 14 day unlimited free trial.](#)

DATADOG **TOOLS** **SPONSORED**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).