

[« Prev](#)[Next »](#)

Database Weekly

Issue 89 — January 29, 2016

FEATURED

19 Resources for Learning About Database Design Errors — A look at numerous online resources for avoiding poor database design that can lead to many problems and cost both time and money.
JEFFREY EDISON

Using AWS Aurora (Amazon's MySQL-Compatible Database) — Advait Shinde shows how his company uses Amazon Aurora to improve scalability by decreasing replica lag and bypassing limitations inherent in enterprise-grade relational databases like AWS RDS.
GOGUARDIAN **VIDEO**

Analyzing A Million Songs on AWS Redshift — A look at extracting data from 1 million songs, uploading it to AWS Redshift via S3, and running some analysis via psql.
MARK LITWINTSCHIK

Unlimited Scalability Unprecedented Speed with the GridGain In-Memory Data Fabric — GridGain enables high-performance transactions, real-time streaming and fast analytics using a distributed, massively parallel architecture which can connect SQL, NoSQL, and Hadoop databases. GridGain is an enterprise-ready solution built on Apache Ignite™.



GRIDGAIN IN-MEMORY DATA FABRIC **SPONSORED**

Postgres Query Plan Visualization — Understanding query plans and EXPLAIN output can be tricky, even for an experienced user. This is a great tool to help better understand how a query is operating.
ALEX TATYANTS **TOOLS**

KeRF: A Next-generation Time-Series and Tick Database — A columnar tick database and time-series language for Linux/OSX/BSD/iOS/Android. Natively speaks JSON and SQL. Ideal for trading platforms, low-latency networking, high-volume analysis of realtime and historical data, etc.
KERF **TOOLS**

Binomial Options Pricing in SQL — How to set up a discrete binomial tree using the binomial options pricing model to value derivatives in SQL, and apply this in an example of pricing a 3-period call option.
PERISCOPE

Gartner Sees Analytics Boom as More Data is Shared — “forecasts greater uptake of emerging analytics technologies like semantic and graph analyses, simulation, complex event processing and neural networks”
DATANAMI

IN BRIEF

[Happy Birthday, Hadoop: Celebrating 10 Years of Improbable Growth](#) — “the first Hadoop cluster went into production at Yahoo 10 years ago”

DATANAM **NEWS**

[Aerospike Adds GeoSpatial Data to Speedy NoSQL](#)

DATANAM **NEWS**

[Get more from Redis: Autoscaling and Analytics from RedisGreen](#) — Redis auto-scaling and analytics.

Backed by support from engineers with years of Redis scaling expertise.

REDISGREEN **TOOLS** **SPONSORED**

[Kafka and Elasticsearch, a Perfect Match](#)

ROLAND KOFLER

[Using Redis for PHP Session Storage](#)

SLAPTIJACK

[Using GraphQL with MongoDB](#)

COMPOSE

[Categorizing Images with Deep Learning Into Elasticsearch](#)

EMMANUEL BENAZERA

[MySQL in Docker or Native – Performance Benchmarks](#)

BALÁZS PÖCZE

[Is Spark Replacing Hadoop?](#)

SD TIMES

[Redis latency](#) — Instantly visible, alertable, and correlated with key eviction rates, memory fragmentation ratio, and more.

DATADOG **TOOLS** **SPONSORED**

[1.7 Billion Reddit Comments Loaded on Google BigQuery](#)

REDDIT

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).