

[« Prev](#)[Next »](#)

Database Weekly

Issue 90 — February 5, 2016

FEATURED

SnappyData: OLTP + OLAP Database built on Apache Spark — A distributed in-memory data store for real-time operational analytics, delivering stream analytics, OLTP (online transaction processing) and OLAP (online analytical processing) in a single integrated cluster. Written in Scala.

SNAPPY DATA, INC [CODE](#)

When Simple SQL Can Be Complex — “A simple statement might have more complexity to it than you might think,” says Daniël, before showing us an example in `SELECT FALSE = FALSE = TRUE;`

DANIËL VAN EEDEN

The Reference Data Pattern: Extensible and Flexible — What if you have dozens and dozens of reference tables? Is there an alternative to the one-table-per-type approach? Learn a generic and extensible database design for handling all your reference data.

JEAN-MARC REYNAUD

How Things Work: Containers and Postgres — The use of containers to simplify and speed the deployment and development of applications is taking off. Most container usage is focused around stateless micro-services, but data and transactions are key components of most applications. Register and attend this informational webinar.

ENTERPRISEDB [SPONSORED](#)



InfluxDB v0.10 GA: Support for 100,000s of Writes Per Second and 98% Better Compression — InfluxDB is an open source time-series database designed for high-perf writes and compact disk storage with built-in clustering.

PAUL DIX [TOOLS](#)

How Graph Databases Stop E-Commerce Fraud in Real Time — Graph databases can be used to detect common patterns in data and behavior surrounding fraud.

NEO4J

Using PostgreSQL in R: A Quick How-To — R is a great language when it comes to analytics, and Postgres is a great database. How about bringing the two together?

NINAZUMEL

BigQuery Under The Hood — A look at Google's large scale data analysis platform. Notably, “it is nearly impossible to starve BigQuery out of resources [...] There is no such thing as one bad query taking down the entire service.”

GOOGLE CLOUD BIG DATA BLOG

JOBS

[Heroku Data Infrastructure Engineer](#) — Join the Heroku Data team and learn about supporting and operating database systems at scale.

HEROKU

[Stop Applying to Jobs - Let Companies Apply to You](#) — On Hired, sign up in 10 minutes and get offers from top companies like Facebook, Uber, & Stripe. Engineers get an average of 5 offers on the platform in 1 week. Try it today.

HIRED.COM

IN BRIEF

[Redis Connector Aims to Boost Spark Performance](#) — An attempt to radically boost Spark performance by using Redis as an in-memory backing store.

DATANAM **NEWS**

[What's New in Elasticsearch 3.0?](#)

JOLICODE **NEWS**

[Redis Performance Dashboard and Autoscaling from RedisGreen](#) — Redis auto-scaling and analytics. Backed by support from engineers with years of Redis scaling expertise.

REDISGREEN **TOOLS** **SPONSORED**

[Diagnosing and Resolving MySQL Deadlocks](#)

BRANDON MAYES

[Summarizing Data in SQL](#)

MATT DELAND

[Running An Elasticsearch Cluster on Docker 1.9 with Swarm and Compose](#)

YOANIS GIL DELGADO

[How To Back Up, Restore, and Migrate PostgreSQL Databases with Barman on CentOS 7](#)

SADEQUL HUSSAIN

[Five Things to Know about Databases that Leverage Partially Homomorphic Encryption](#)

DAVID CROOKER

[If You're A DB, You Need to Manage CPU, Memory and IO](#)

OREN EINI

[Hadoop's Second Decade: Where Do We Go From Here?](#)

DATANAM

[SQLDust: Complex SQL Generation from Elixir](#)

PAUL ENGEL **CODE**

[Solve SIEBEL OM Crashes](#) — Automated Siebel CRM Repository Review to solve OM crashes, data corruption, slow transaction, workflow and more for \$1K in under 4 hours.

GERMAIN SOFTWARE **TOOLS** **SPONSORED**

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).