Database Weekly

Issue 98 — April 1, 2016

FEATURED

Elastic Gives ElasticSearch a Graph Query Engine — Elastic has announced that it's added a graph query engine to Elasticsearch, so users now have the option of using their search indexes as the basis for conducting graph analyses.
DATANAMI

Caravel: Airbnb's Data Exploration Platform — A data exploration platform designed to be visual, intuitive and interactive. The video and screenshots are very appealing. Caravel allows visualization-based exploration of data while performing fast "slicing and dicing" against any dataset.
MAXIME BEAUCHEMIN `TOOLS`

How Neo4j Saved my Relationship — A look at the power of graph databases in taming complex data problems.
CORALINE ADA EHMKE `VIDEO`

Get more from Redis: Autoscaling and Performance Dashboard from RedisGreen — Ever wonder what your Redis server was actually doing? Find out with RedisGreen's dashboard and monitoring tools.
REDISGREEN `SPONSORED`

Citus Unforks From PostgreSQL, Goes Open Source — Citus Data has released Citus 5.0 which comes along with two major announcements. The first is it's entirely an extension, making it the first distributed databased to not fork Postgres but use native extension APIs. The second, it's now open source.
CITUS DATA

SQL: How to Create An Execution Plan — When you're trying to get SQL to run faster, there's one thing that's vital: an execution plan. In this post Chris Saxon investigates how to create yours.
ORACLE

What is Blockchain? — A primer for people just starting to wrap their heads around blockchain, a type of distributed database that maintains an expanding ledger of data and records.
W. IAN O'BYRNE

Analyzing A Billion Taxi Rides in PostgreSQL — We covered when Mark analyzed a billion rides in Redshift, now here's the same but using Postgres and cstore, the columnar storage FDW from Citus Data
MARK LITWINTSCHIK

JOBS

IN BRIEF

SQL Server Developer Edition Is Now Included for Free with Visual Studio Dev Essentials
MICROSOFT `NEWS`

Gartner Magic Quadrant for Data Warehouse Solutions — This report will help data, analytics and IT leaders find the right vendor for their needs.
MEMSQL `NEWS` `SPONSORED`

CockroachDB Skitters into Beta — First unveiled last June, CockroachDB has now entered beta.
SPENCER KIMBALL `NEWS`

Machine Learning Templates with SQL Server 2016 R Services — In-database analytics in SQL Server 2016 allows users to combine the power of SQL Server and Microsoft R Server (or Open Source R), without data leaving the database.
MICROSOFT `NEWS`

SQLite 3.12.0 Released
SQLITE `NEWS`

Learn to Live with Apache Hive in 12 Easy Steps — Hive lets you use SQL on Hadoop, but tuning SQL on a distributed system is different. Here's 12 tips to help.
ANDREW C. OLIVER

Five Ways to Paginate in Postgres, From The Basic to The Exotic — This article examines several methods of server-side pagination and discusses their tradeoffs when implemented in PostgreSQL.
JOE NELSON

How to Validate User Input in A NoSQL Web Application — Practical approaches on how to validate user input for a schema-less database.
RYAN PAUL

PipelineDB Persistence with Flocker and Docker Swarm — Covers running a basic PipelineDB server container on a host machine using Docker.
STEPHEN NGUYEN

A Simple Guide To Building Reports Using SQL Server Reporting Services — A detailed outline of the steps from creating the stored procedure that will run the report to building a front-end report using the Business Intelligence module in Visual Studio.
JACK WORTHEN

Why You Should (Sometimes) Use MongoDB
NICHOLAS JOHNSON

To SQL or NoSQL? That's The Database Question — Lisa Vaas opines that as the technological lines blur, there's not always a clear-cut answer between the two.
ARS TECHNICA

[Picking The Right Caching Tier for Your Database](#) — Caching tiers are used to lower latencies under high database throughput, so it pays to pick the right approach.
CIHAN B

[How NoSQL Graph Databases Still Usurp Relational Dynasties](#)
ANDREW COBLEY

[Unlimited Scalability & Unprecedented Speed with In-Memory Computing](#) — GridGain enables high-performance transactions, real-time streaming and fast analytics using a distributed, massively parallel architecture.
GRIDGAIN IN-MEMORY DATA FABRIC `TOOLS` SPONSORED

[Mrs. Null: These Unlucky People Have Names That Break Computers](#) — To any programmer, it's painfully easy to see why "Null" could cause problems for software interacting with a database.
CHRIS BARANIUK