DB Weekly

Database Weekly

FEATURED

Redis 3.2.0 Released with Lots of New Features — Redis, an in-memory, high performance 'data structure server', is a key piece of many apps now. 3.2 introduces a new 'GEO' API for indexing lat/long points, a BITFIELD command, and more.
SALVATORE SANFILIPPO

Spark 2.0 Technical Preview: Easier, Faster, and Smarter — Spark 2.0 (the large scale data processing engine) has greatly improved SQL support and faster query times. The final release is a few weeks away but you can get early access to its features now.
REYNOLD XIN

Discussing SQLite with Richard Hipp — An hour long talk with Richard Hipp, creator of SQLite, the most widely deployed DB engine in the world.
THE CHANGELOG

Connecting Hadoop and EDB Postgres to Shrink Big Data Challenges — By combining data stored in Hadoop clusters with its transactional information, EDB and Hortonworks provide a seamless view of information that will empower organizations to engage in a deeper analysis of corporate data and find new efficiencies for customer-driven revenue programs.
ENTERPRISEDB SPONSORED

An Open Source Tour de Force at Apache: Big Data 2016 — A look at how open source projects are driving big data, particularly those being incubated by Apache.
DATANAMI

Realm Xamarin: A New Reactive Database for .NET Developers — A mobile DB built specifically for Xamarin that offers easy object persistence and full query capabilities.
REALM TOOLS

PostgreSQL Scalability: Towards 1 Million TPS — While Postgres becomes more usable with every release, it's also made big performance improvements. Take a look at how close Postgres is to hitting one million transactions per second.
ALEXANDER KOROTKOV

JOBS

IN BRIEF

'Great to see that my name still causes SQL errors..'

TWITTER `NEWS`

Spark 2.0 Will Be Faster, Easier for App Development, and Tackle Streaming Data

LINUX.COM `NEWS`

Redis Gaining a Loadable Module System — The potential use cases for Redis could explode from this.

SALVATORE SANFILIPPO `NEWS`

On-demand Webinar: Turbocharge Your SQL Queries In-Memory with Apache® Ignite™ — Learn how to boost performance 1,000X and scale to over 1 billion transactions per second with your SQL-based applications.

GRIDGAIN IN-MEMORY DATA FABRIC `SPONSORED`

How to Create a Redis Module in 15 Lines of Code — Based on the news that Redis is gaining a module system (above).

DVIR VOLK

A Quick Guide to Redis 3.2's Geo Support

COMPOSE

The Basics of Database Transactions

KUMAR RAJEEV RASTOGI

Seldom Used Postgres Functions: bool_or and bool_and

RIMAS SILKAITIS

From SQL to Microservices: Integrating AWS Lambda with Relational Databases

AMAZON

SQLUnplugged Chat About SQL Server 2016 GA and SQL on Linux

CHANNEL 9 `VIDEO`

Bridging NoSQL Databases Into Open Data Science Initiatives

IBM BIG DATA AND ANALYTICS HUB

Stop building APIs — DreamFactory auto-generates APIs from nearly any database, instantly. Out of the box support for V8 (JS), Node.js, PHP, and Python for custom scripting.

DREAMFACTORY `TOOLS` `SPONSORED`

6 Unsung Apache Big Data Projects That Aren't Spark

LINUX.COM `TOOLS`

Teleport: A Trigger-Based Postgres Replicator Written in Go — A trigger-based replicator that performs DDL migrations by diffing schema changes and replicating real-time data changes based on DML triggers. No special permissions required, so it's suitable for AWS RDS.

PAGAR.ME `TOOLS`

membership.db: SQL Boilerplate for User Accounts, Roles & Auth Tokens for Webapp Use

MEMBERSHIP DATABASE STARTER KIT `CODE`

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address is safe — here's our privacy policy.