

[« Prev](#)[Next »](#)

Database Weekly

Issue 106 — May 26, 2016

FEATURED

Apache Spark as a Compiler: Joining a Billion Rows on your Laptop — A look at how Apache Spark acts as a compiler to optimize queries using whole-stage code generation techniques commonly used in language compilers.

SAMEER AGARWAL

EnterpriseDB Wraps PostgreSQL Into An Enterprise-Grade Suite to Challenge Oracle — EnterpriseDB has brought more tools into its Postgres distro to challenge big commercial DBs.

DANIEL ROBINSON

A Look at the Design of SQLite4 — Described as “just like SQLite3, but with an improved interface and file format”, SQLite4 is going to be “an alternative, not a replacement, for SQLite3.”

SQLITE

TrailDB: An Efficient Library for Storing and Processing Event Data — AdRoll has open sourced a fast, C-based library for compressing and handling time-series event data.

VILLE TUULOS [CODE](#)

Have you got SQL fingers? — Watch these free SQL Prompt tips videos for SQL writing hints from top SQL Server MVPs. SQL Prompt is the SQL code productivity add-in for SQL Server Management Studio and Visual Studio. Find out how easily you can write SQL.

RED GATE [SPONSORED](#)



Unicorn: BigTable, Document and Graph Database with Full Text Search — A simple abstraction of BigTable-like databases such as Cassandra, HBase, and Accumulo that provides an easy-to-use document data model and MongoDB-like API.

HAIFENG LI [TOOLS](#)

Benchmarking Elasticsearch vs InfluxDB for Time-Series Data — Some of the InfluxData team set out to compare the performance and features of InfluxDB and Elasticsearch for common time-series workloads. (*Note: InfluxData develops InfluxDB.*)

SHUBHRAKAR [VIDEO](#)

Understanding Outliers with Skew and Kurtosis in SQL — *Skewness* is a measure of the asymmetry of the probability distribution of a real-valued random variable about its mean, and kurtosis is a measure of the “tailedness” of the same distribution. Confused? Here’s a practical example.

PERISCOPE

Efficient Pagination in SQL and ElasticSearch — Traditional pagination queries can have a high hidden cost. Here, an interesting technique to make pagination more efficient is explored.

SALSIFY

Monitoring MongoDB Performance Metrics — By properly monitoring MongoDB you can quickly spot slowdowns, or pressing resource limitations, and know how to fix performance issues. This series covers WiredTiger but there's a [MMAP edition](#) too.

JEAN-MATHIEU SAPONARO

Scaling to 100M: MySQL is a Better NoSQL — When considering a NoSQL use case, such as key/value storage, MySQL made more sense to Wix in terms of performance, ease of use, and stability.

YOAV ABRAHAM

JOBS

Job Offers. No resume necessary. — Create your Hired profile to get top companies to start applying to hire you. Get offers from \$75,000 - \$250,000 on the platform in 1 week.

HIRED.COM

IN BRIEF

[How Spark and Hadoop Are Advancing Cancer Research](#)

ALEX WOODIE **NEWS**

[Here's The Slide Deck MySQL Used to Raise \\$16 Million in 2003](#)

ROBIN WAUTERS **NEWS**

[Amazon Redshift Gets Up to 2X Throughput and 10X Vacuuming Performance Improvements](#)

AMAZON WEB SERVICES **NEWS**

[Cray Unveils Open Source 'Big Data' Box](#) — Serious metal with 35TB of PCIe SSD storage, 22TB of RAM, and preloaded with OpenStack and Mesos.

THE REGISTER **NEWS**

[How we improved query response times 6X with Apache Storm and Cassandra](#) — Keen IO platform engineer shares lessons learned while scaling out their analytics infrastructure with Apache Storm and Cassandra.

KEEN IO **SPONSORED**

[Data Types in PostgreSQL](#) — A 128 page slidedeck 'guided tour' through PostgreSQL data types.

PETER VAN HARDENBERG

[Postgres & Redis Sitting In a Tree](#) — “Typically a company may start with a relational database like Postgres and then add Redis for more high velocity use-cases. What if you could tie the two systems together to enable so much more?”

RIMAS SILKAITIS

[Converting MySQL Syntax and Functions into Redshift](#)

PERISCOPE

[Event Logging with ElasticSearch from the Browser](#)

EDGAR RIBEIRO **TOOLS**

[MySQL By Examples for Beginners](#)

CHUA HOCK-CHUAN

[Creating Pivot Tables in PostgreSQL Using the Crosstab Function](#) — The 'tablefunc' extension provides a really interesting set of functions. One of them is the crosstab function, used for pivot table creation. That's what this article covers.

MARIAALCARAZ

[How to Move From Amazon RDS to A Dedicated PostgreSQL Server](#)

AUTHENTISE, INC

[Is A RID Lookup Faster Than A Key Lookup On SQL Server?](#)

AARON BERTRAND

[7 Reasons to Choose A File System \(SolFS\) Over a DB for Managing App Data](#)

ELDOS

[DreamFactory 2.1.2 released, includes app packages](#) — Join data between a SQL database and MongoDB with a few configuration clicks.

DREAMFACTORY **TOOLS** **SPONSORED**

[SQL and NoSQL Databases Ranked by StackOverflow Activity](#) — MySQL and SQL Server are well out in the lead.

STACKOVERKILL **TOOLS**

[countminsketch: An Apporximate Frequency Counter Module for Redis](#)

REDIS LABS MODULES **CODE**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).