DB Weekly

Database Weekly

Issue 122 — September 16, 2016

FEATURED

MySQL 8.0.0 Milestone Release Available — A development milestone release so you can play with MySQL 8's new features now.
MYSQL SERVER TEAM

HTSQL: A Navigational Query Language for Relational DBs — Designed for data analysts and 'accidental programmers' who have complex business inquiries to solve and need a productive tool to write and share database queries. Open source.
PROMETHEUS RESEARCH TOOLS

Neo4j + KeyLines: The Developer's Route out of Dependency Hell — How to use Neo4j in conjunction with KeyLines for software dependency management using insights gleaned from graph visualization.
MIRO MARCHI

Auto-generated REST APIs from any database — Version 2.3 of DreamFactory's powerful open source API automation platform brings major enhancements to speed, caching, and versatility. Free demo.
DREAMFACTORY SPONSORED

BlazingDB Uses GPUs to Manipulate Huge Databases in No Time — BlazingDB works with GPU instances on popular cloud services to let you rapidly run SQL queries over large amounts of data.
ROMAIN DILLET

MIT Programmers Attack Big Data Memory Gap — A new programming language that efficiently manages memory during big data operations.
DATANAMI

Next Generation Tools for Data Science: Apache vs Google — A comparison of Apache Spark and Google Cloud Dataflow in the context of being two key 'next generation' data processing frameworks.
UNOFFICIAL GOOGLE DATA SCIENCE BLOG

QuineDB: A Quine That Is Also A Key-Value Store — A quine is a program that, when run, outputs its own source code. Stores data in a shell variable. Solely for fun and an interesting experiment.
GARY FREDERICKS CODE

JOBS

IN BRIEF

Some Notes on Shipping CouchDB 2.0 — CouchDB is "tantalisingly close to shipping 2.0".
JAN LEHNARDT `NEWS`

SIEUFERD: A General Purpose UI for Relational Databases — An academic project that's not publicly available.
EIRIK BAKKE `NEWS`

Unraveling Hadoop and Spark Performance Mysteries
DATANAMI `NEWS`

Calculating Proportional Values in SQL — Aggregations are common and generally straightforward in SQL, but some more complex operations like medians require a bit more work as demonstrated here.
PERISCOPE

MySQL 8.0's Invisible Indexes — Invisible indexes are a new feature in MySQL 8.0 to hide an index from MySQL's optimizer.
MARTIN HANSSON

An Overview of MySQL's Information Schema
EMIL DRKUŠIĆ

Recovering a Postgres Database From Disk Level Corruption
DENISH PATEL

Neo4j With Scala: Migrating Data From Other Databases to Neo4j
ANURAG KNOLDUS

Fun with SQL: Computing Run Rate and Month Over Month Growth
CITUS DATA

Command-Line Tools Can Be 235x Faster Than Your Hadoop Cluster — A 'golden oldie' that's done the rounds again this week.
ADAM DRAKE

Building a Recommendation Engine with Machine Learning Techniques
BRIAN SAM-BODDEN `VIDEO`

Avoid Using COUNT() in SQL When You Could Use EXISTS()
JOOQ

Why You Should Build Your Own NoSQL Database — Covers the basic principles and creates something very simple in 33 lines of Crystal.
MARCELO BOEIRA

5 Mistakes to Avoid When Deploying An Analytical Database — Leading data-driven organizations point out five common pitfalls.
ALICE LAPLANTE

[Spot performance bottlenecks right away with Datadog](#) — Correlate and investigate: Monitor query latency and errors alongside CPU, memory, network, and I/O metrics.

DATADOG `TOOLS` `SPONSORED`

[sQucumber-redshift: A Framework for Defining and Executing SQL Integration & Acceptance Tests on AWS Redshift](#)

STEFANIE GRUNWALD `TOOLS`

[EliasDB: A No Dependencies Graph Database Written in Go](#) — Exposes a REST API.

MATTHIAS LADKAU `CODE`

[« Prev](#)

[Next »](#)