

[« Prev](#)[Next »](#)

Database Weekly

Issue 125 — October 7, 2016

FEATURED

Data Denormalization Is Broken — Why it's hard to write good app-layer code for everyday business logic. *"I think the database is overdue for a paradigm shift in the way it handles denormalization. I predict powerful denormalization engines will be invented in the next few years."*

LIRON SHAPIRA

Neural Redis: Neural Networks Module for Redis — A module that implements feed-forward neural networks as a native data type in Redis. The goal is to provide a simple way for anyone to get started with machine learning.

SALVATORE SANFILIPPO [CODE](#)

Why You Should Design Your Database to Optimize for Statistics — Databases are increasingly being used to collect together aggregate data, histograms and overall statistics. Keep this in mind when designing things.

JOQQ

How can you help your team write better, shareable SQL faster? — Learn more with our free video tips that show how SQL Prompt lets you write, refactor, and share SQL. An add-in for SQL Server Management Studio and Visual Studio, it strips away the repetition of coding and standardizes it everywhere.



RED GATE [SPONSORED](#)

RethinkDB Is Shutting Down — The business behind RethinkDB, an innovative open source realtime database, is closing down and its employees moving to Stripe. The project will live on as open source, however.

RETHINKDB

Lessons From MySQL Database Failure — Backups, verification, replication, and database security are covered in this 31-minute talk looking at lessons learnt from MySQL failures at scale.

COLIN CHARLES [VIDEO](#)

Accelerating Slow Databases That Wear People Down — A look at how a frustration with database speeds inspired the creator of MapD and how MapD is using various techniques to boost database performance.

TIMOTHY PRICKETT MORGAN

The Third Age of Data — The massive amounts of data being generated hold the promise of ever-greater insight, but only for those who successfully process and harness the flood of information.

JEFF COBB

Debugging PostgreSQL Performance, The Hard Way — When JustWatch saw a disproportionate increase in database reads they decided to find out why — instead of just throwing money at the problem.

JUSTWATCH

[Don't Get Frustrated Job Hunting - Get Hired](#) — Sick of pushy recruiters, and dead end interviews? Try Hired to hear from top tier companies, and only talk to relevant companies.
HIRED

IN BRIEF

[RethinkDB Needs A New Home](#) — A form looking for feedback on how RethinkDB should be handled by the open source community. *Note: This form is not managed by the RethinkDB team, but by open source contributors.*

GOOGLE.COM **NEWS**

[Introducing MariaDB MaxScale 2.0](#) — MaxScale is a database proxy that sits between the application and database layer.

MARIADB **NEWS**

[MySQL 8.0 Labs: JSON Aggregation Functions](#)

CATALIN BESLEAGA **NEWS**

[ScyllaDB now on Compose](#) — Blazing fast replacement for Cassandra that boasts 1M transactions sec/node. Get ScyllaDB started with your favourite runtimes.

COMPOSE **SPONSORED**

[How to Solve 5 Elasticsearch Performance and Scaling Problems](#) — Five common Elasticsearch challenges, and how to deal with them.

EMLY CHANG

[Working with SQLite Databases using Python and Pandas](#)

VIK PARUCHURI

[Pairing LLVM JIT with PostgreSQL Can Speed Up Database Performance](#) — LLVM JIT is used for just-in-time compilation of queries.

DMTRY MELNIK **VIDEO**

[How MongoDB Helps Telefonica Digital Accelerate Time to Market](#) — Why telecommunications giant Telefonica migrated from Oracle to MongoDB. *37 minutes.*

MONGODB **VIDEO**

[Yandex Mail's Oracle to Postgres Success Story](#)

YOUTUBE **VIDEO**

[Has Big Data Peaked?](#) — Fewer companies are planning to invest in big data projects over the next two years due to concerns over ROI.

ALEX WOODIE

[6 Reasons to Choose PostgreSQL 9.6](#)

DAMEN CLOCHARD

[When A Disk Cache Performs Better Than An In-Memory Cache \(Befriending The .NET GC\)](#)

PRODUCTIVE DAN

[Get real-time visibility into dynamic environments with Datadog](#) — One-line installation. Two-minute configuration. Total infrastructure observability is a couple clicks away.

DATADOG **SPONSORED**

[Mirage: An Online GUI for Writing Elasticsearch Queries](#)

APPBASEIO.GITHUB.IO [TOOLS](#)

[elasticsearch_loader: Batch Load Data Files into Elasticsearch](#)

MOSHE ZADA [TOOLS](#)

[ZSON: A PostgreSQL Extension for Transparent JSONB Compression](#)

ALEKSANDER ALEKSEEV [CODE](#)

[pgmigrate: Simple Tool to Evolve PostgreSQL Schemas Easily](#) — Yandex just completed one of the larger Postgres migrations from MySQL and have now open sourced their tool to make it easier to evolve.

YANDEX [CODE](#)

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).