Database Weekly

Issue 142 — February 17, 2017

FEATURED

Google Unveils Cloud Spanner: A Global Database Service — A globally-distributed relational database service with ACID transactions and SQL semantics. WIRED's writeup gives the bigger picture of why this is a big deal.
GOOGLE

Is PostgreSQL Good Enough (For Almost Everything)? — Can Postgres be 'good enough' for all sorts of different use cases? Time series, object stores, pubsub, realtime, and more are covered here.
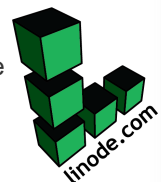RENE DUDFIELD

Tarantool: The Good, The Bad and The Ugly — A practical demonstration of Tarantool, an open source in-memory NoSQL database and Lua app server.
VADIM POPOV

Linode is the SPEEDIEST SSD host for your DB - simple, reliable & powerful. — MySQL? MariaDB? PostgreSQL? Whatever DB you use, run it on a Linode, the most scalable, reliable and fastest servers in the cloud. Use promo code `DB20` for a $20 credit and get started.
LINODE <mark>SPONSORED</mark>

Microsoft Azure DocumentDB Goes After MongoDB Users — Microsoft is encouraging MongoDB users to switch to their Azure DocumentDB service, promising no rewrites would be necessary due to its MongoDB protocol support.
MICROSOFT

A Look at Instapaper's Outage Cause and Recovery Process — Bookmarking service Instapaper experienced a 30 hour outage. Here's the story of how their MySQL database failed, why, and how it was resolved.
BRIAN DONOHUE

'Climbing' Big Data With Virtual Reality — Why just look at graphs on a screen when you can go into virtual reality and 'climb' your data? Video here.
DATANAMI

PostgreSQL At 10TB and Beyond — Chris Travers discusses what happens when managing over 10 terabytes of data in PostgreSQL. Does it scale and what kinds of problems need to be resolved?
CHRIS TRAVERS <mark>VIDEO</mark>

IN BRIEF

[Google Launches Video Analytics Challenge](#) — Essentially it involves classifying YouTube videos. Prizes up for grabs.

DATANAMI

[Citus 6.1 Released: Horizontal Scaling for PostgreSQL](#)

CITUS DATA

[Inside IBM ML: Real-Time Analytics On the Mainframe](#)

DATANAMI

[Online.net Introduces ODS: High Performance SQL Database Infrastructure](#) — A fully managed cloud SQL service.

ONLINE.NET

[Practical Query Optimization for MySQL and PostgreSQL](#) — In this book, you'll learn how to find, evaluate, and improve the queries running on your database servers.

VIVIDCORTEX

[Filling Missing Data and Plugging Gaps by Generating a Continuous Series in SQL](#)

SILOTA

[Pipelines vs. MapReduce to Speed Up Data Aggregation in MongoDB](#)

HENRY MOLLMAN

[Managing Houses and Properties: A Real Estate Agency Data Model](#) — The latest data model dive from Vertabelo.

EMIL DRKUŠIĆ

[The Performance Impact of SQL's 'WITH'](#)

MARKUS WINAND

[How Yahoo's Internal Hadoop Cluster Does Double-Duty on Deep Learning](#)

NICOLE HEMSOTH

[Fast and Flexible Query Analysis at MapD, with Apache Calcite](#) — Yet another [Apache database project](#) :-)

ALEX SUHAN

[UUID or GUID as Primary Keys? Be Careful](#)

TOM HARRISON JR

[Why Google's Spanner Database Won't Do As Well As Its Clone](#)

TIMOTHY PRICKETT MORGAN

[Life After 1 Year of Using Neo4j](#)

ANDREY NIKISHAEV

[RethinkDB Versus PostgreSQL: My Personal Experience](#) — A personal comparison only: *"It's what I would tell you if we were talking by the water cooler."*

WILLIAM STEIN

[Postgres Worst Practices (A Slidedeck)](#)

ILYA KOSMODEMIANSKY

[TensorFlowOnSpark: Brings TensorFlow Programs Onto Apache Spark Clusters](#)

GITHUB.COM `TOOLS`

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address is safe — here's our privacy policy.