

[« Prev](#)[Next »](#)

Database Weekly

Issue 146 — March 17, 2017

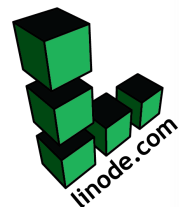
FEATURED

How Discord Indexes Billions of Messages — Discord is a chat service supporting millions of users. How could they add quick searching? Here's how they set it up with Elasticsearch.
DISCORD

Timescale: An Open Source Time-Series Database — SQL made scalable for time-series data. It's Postgres compatible and optimized for fast ingest & complex queries.
TIMESCALE

Real-Time Deduping At Scale with Redis — Deduping at scale is a hard problem and Tapjoy share a lot of details of how they use Redis to do it on over 2 million analytics messages per minute arriving in an 'at-least-once' style.
TAPJOYENGINEERING

Linode is the SPEEDIEST SSD host for your DB - simple, reliable & powerful. — MySQL? MariaDB? PostgreSQL? Whatever DB you use, run it on a Linode, the most scalable, reliable and fastest servers in the cloud. Use promo code DB20 for a \$20 credit and get started!
LINODE **SPONSORED**



Miller: Like awk, sed, cut, join, and sort for CSV, TSV, and JSON — Provides Unix tool-esque functionality but around key-value pair data. Written in C too, so no dependencies.
JOHN KERL **TOOLS**

Hoodie: Uber's Incremental Processing Framework on Hadoop — Uber Engineering recently built and open sourced Hoodie, an incremental processing framework to power Uber's data pipelines at low latency.
UBER ENGINEERING **TOOLS**

Has Hadoop Failed Us? — Alex Woodie brings together several people from across the industry to question if the Hadoop dream of unifying data and compute in a distributed manner has become too expensive and complex.
DATANAM

Don't Migrate Databases Automatically — Migrating your DB and deploying app code at the same time can create a race condition that could take down your app.
PHILIP I. THOMAS

IN BRIEF

[Amazon Aurora Cuts Entry-Level Pricing With Support for t2.small Instances](#)

AMAZON WEB SERVICES, INC. **NEWS**

[FaunaDB Launches Serverless Cloud](#)

FAUNA **NEWS**

[Database Performance Monitoring Buyer's Guide](#) — This guide is designed to aid when evaluating database monitoring solutions for your unique environment.

VIMDCORTEX **SPONSORED**

[Faster SQL Through Occasionally Choosing Natural Keys Over Surrogate Keys](#)

JAVAOOQ

[Amazon Redshift Pitfalls And How To Avoid Them](#)

JAMES KATZ

[The Complete Guide to GROUP BY on Oracle](#)

BEN BRUMM

[An Antique Store Data Model](#) — Vertabelo's fine data modelling guides continue to get more niche :-)

VERTABELO

[First Impressions of Google Cloud SQL for PostgreSQL](#)

PAYAL SINGH

[Why You Should Choose SQL for Your Web Application](#)

PARIS KASIDIARIS

[SQL Source Control: track each change to your SQL Server database](#) — Get a full history in your source control system. See who made changes, what they did and why. [See how](#).

RED GATE **TOOLS** **SPONSORED**

[Euphrates: A Periodic 'Table Copier' from MySQL to Amazon Redshift](#)

PATREON **TOOLS**

[Flyway: An Open Source Database Migrations Tool](#) — Works with all major databases.

BOXFUSE **TOOLS**

[TeamSQL: A Multi-Platform SQL Client \(in Private Beta\)](#) — Looks promising. You can join a waiting list if interested.

TAPHOLIC **TOOLS**

[A Curated Guide to the Best Tools and Resources for Data Visualization](#)

DATAMZ.TOOLS **TOOLS**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).