

[« Prev](#)[Next »](#)

Database Weekly

Issue 149 — April 7, 2017

FEATURED

Building A Scalable Time-Series Database on Postgres — [TimescaleDB](#) is a new open-source time-series database optimized for fast ingest and complex queries.

AJAY KULKARNI

SQLite 3.18 Released — `PRAGMA optimize` has been added, a mechanism for attempting to optimize the database and improve long-term query performance.

SQLITE

Apache Quickstep: A SQL-Based Performance-Focused DB Engine — Originating in 2011, but now being incubated at Apache, Quickstep competitively runs SSB and TPC-H benchmarks on a single node and is designed to support the latest in hardware technology for faster performance. [GitHub repo](#).

APACHE FOUNDATION **TOOLS**

Webinar: Better Azure SQL Performance — Learn how to reduce development times by getting a better understanding of how your Azure SQL workloads are really performing. [Register now](#).

SELECTSTAR **SPONSORED**



A Data Model for an Electric Power Distribution System — Vertabelo's database modelling scenarios continue to get more niche but it's always interesting to see how a professional database designer approaches the challenges involved.

VERTABELO

Business Intelligence Performance Benchmarks with Google BigQuery — A deeper look at Google's BigQuery, and how it stacks up in the BI-on-Big Data ecosystem.

JOSHUA KLAHR

Redis Flash: Double Performance on NVMe SSDs — Redis Flash on [AWS I3 instances](#) (which have NVMe SSDs) gets up to 2.5x performance compared to on I2s.

KEREN OUAKNINE

IN BRIEF

[Meet Ray, The Real-Time Machine-Learning Replacement for Spark](#)

ALEX WOODIE **NEWS**

[Data Compression Improvements in Amazon Redshift Bring Compression Ratios Up to 4x](#)

AMAZON **NEWS**

[Getting Started with Apache Cassandra \(with Ruby\)](#)

VITO BOTTA **TUTORIAL**

[Language-Specific Views in MongoDB 3.4](#) — Building multiple filtered, language-specific views of a common collection using collations & partial indexes.

MONGODB **TUTORIAL** **SPONSORED**

[Docker: Getting Started with Python, Redis, and NGINX](#)

ROMAN GAPONOV **TUTORIAL**

[Record Linking with Apache Spark's MLlib and GraphX](#) — A scalable approach to fuzzy data matching.

TOM LOUS **TUTORIAL**

[Distributed COUNT\(DISTINCT\) with HyperLogLog on Postgres](#)

CITUS DATA **TUTORIAL**

[Deploy a Couchbase + Node Webapp as Docker Containers](#)

NIC RABOY **TUTORIAL**

[Backing Up SQL Server Databases Hosted on Azure VMs](#)

DATABASE JOURNAL **TUTORIAL**

[Why is January 1, 1753 The Minimum Date Value in SQL Server?](#) — It's Philip Stanhope, 4th Earl of Chesterfield's fault.

STACK OVERFLOW **STORY**

[How Can I Impress at a DBA Job Interview With Little Experience?](#)

REDDIT **OPINION**

[How to find and fix the slowest code in your .NET application](#) — Find bottlenecks in your code or database & boost performance with ANTS Performance Profiler. [Try it free now.](#)

RED GATE **TOOLS** **SPONSORED**

[Ifrit: An SQL to MongoDB Query Compiler for Data Aggregation](#) — Written in PureScript, a language that compiles to JS.

MATTHIAS BENKORT **TOOLS**

[Prophet: Forecasting At Scale from Python and R](#)

FACEBOOK RESEARCH **TOOLS**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).