

[« Prev](#)[Next »](#)

## Database Weekly

Issue 151 — April 21, 2017

## FEATURED

**GeoMesa: An Open-Source Spatio-Temporal Database Layer** — GeoMesa provides spatio-temporal indexing on top of Accumulo, Bigtable, and Cassandra, as well as near real-time stream processing and spatial semantics on top of Kafka.

THE GEOMESA PROJECT **TOOLS**

**Amazon Redshift Spectrum: Exabyte-Scale In-Place Queries of S3 Data** — Spectrum makes it possible to run complex queries using Redshift on data stored on AWS S3 without any loading or data prep.

AMAZON **TOOLS**

**The Top Features Coming to SQL Server 2017** — A summary of what Microsoft has unveiled as coming to SQL Server 2017 from [Python support](#) to adaptive query optimization and [a built-in graph database](#).

JOEY D'ANTONI

**High Available and Scalable Open Source Database - SiriDB** — The time series database SiriDB can scale on the fly, is robust by design and uses a unique mechanism to operate without indexes. SiriDB's query language includes dynamic grouping for easy and fast analysis. [GitHub repo](#)

TRANSEPTOR TECHNOLOGY **SPONSORED**

**How to Calculate Multiple Aggregate Functions in a Single Query** — A look at alternatives to using a large collection of SQL queries to query the same data in different ways, such as pivots and grouping sets.

JAVAOOQ

**Architecture of Giants: Data Stacks at Facebook, Netflix, Airbnb, and Pinterest** — Simple event data infrastructure diagrams from several fast-scaling companies.

MICHELLE WETZLER

**Why to Use A Relational Database for Time-Series Data** — NoSQL databases are commonly used to store time-series data, but the creator of [TimescaleDB](#) sets out a technical case for bringing time-series data into a relational setup.

MIKE FREEDMAN (TIMESCALE)

## IN BRIEF

## [Iguazio Re-Architects the Stack for Continuous Analytics](#)

DATANAM **NEWS**

[Oracle's Databases and Developer Tools Now Available on Docker](#) — Oracle's main database product, WebLogic, Coherence, MySQL, and others are available in Docker containers on the Docker Store marketplace.

DOCKER **NEWS**

[Whitepaper: 5 Steps to Agile DB Management](#) — Database Management is years behind software development. Here's how to bring your DBs up to speed.

SELECTSTAR **SPONSORED**

[Removing Duplicate Rows in Postgres](#) — What if you accidentally load data twice? A look at handling duplicate data and a resulting clean up.

HANS-JUERGEN SCHOENIG **TUTORIAL**

[Tips for Monitoring Redis](#) — Ways to get more info from Redis, such as on latency and slow commands.

MIKE PERHAM **TUTORIAL**

## [Creating Role-Based Access Control in MongoDB](#)

SCALEGRID **TUTORIAL**

## [A Data Model for the Recruitment Process](#)

VERTABELO **TUTORIAL**

## [MySQL Stored Procedures 101](#)

PETER LAFFERTY **TUTORIAL**

## [Inside Anodot's Anomaly Detection System for Time-Series Data](#)

DATANAM **STORY**

## [Designing a Data Warehouse on Hadoop](#)

REALITY GAMES **STORY**

[Running Out of IDs](#) — *"Keep on using serial for most use cases and keep bigserial in your back pocket if a real need arises."*

JOSH BRANCHAUD **STORY**

[Cassandra vs. MongoDB](#) — Considering the differences between Cassandra vs. MongoDB as the data store for your next project

SCALEGRID.IO **OPINION** **SPONSORED**

## [BigQuery: One Store to Rule Them All](#)

BRIGADE ENGINEERING **OPINION**

## [SQL Is The Perfect Interface](#)

DAN GOLDIN **OPINION**

## [Apache ORC: High-Performance Columnar Storage for Hadoop](#)

APACHE FOUNDATION **TOOLS**

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).