

[« Prev](#)[Next »](#)

Database Weekly

Issue 155 — May 19, 2017

FEATURED

Google's 'Cloud Spanner' Is Now Production Ready — A globally-distributed relational database service with ACID transactions and SQL semantics. [WIRED's writeup](#) (from February) gives the bigger picture of why it's a big deal.

GOOGLE CLOUD PLATFORM BLOG [TOOLS](#)

Google BigQuery Now Offers 10GB of Free Storage — BigQuery has long provided an easy way to analyze massively large datasets and now there's a free storage tier to make getting started even easier.

GOOGLE CLOUD PLATFORM

PostgreSQL 10 Beta 1 Released — A big release is on the way later this year featuring logical replication and native table partitioning. Here's [slides covering the changes](#), plus a look at [incompatible changes](#).

POSTGRESQL

Whitepaper: 5 Steps to Agile Database Management — Database Management is years behind software development. Here's how to bring your DBs up to speed.

SELECTSTAR [SPONSORED](#)

Google Spanner: Becoming a SQL System — An academic paper presented at SIGMOD 17 that digs into the specifics of the *Spanner* globally-distributed data management system Google uses (as mentioned above).

ACM/ GOOGLE

Badger: A High-Perf Key/Value Storage Engine — An embeddable, persistent, simple and fast key-value store, written natively in Go. [GitHub repo](#).

DGRAPH [CODE](#)

Neo4j 3.2 Released with Multi Datacentre Support — It's now easier to run the graph database across continents. There are also production governance (e.g. Kerberos auth), query monitoring, and performance improvements.

NEO4J

IN BRIEF

[Committers Talk Hadoop 3 at Apache Big Data](#) — Focusing mostly on the forthcoming addition of erasure coding to HDFS which is well explained here.

DATANAM **NEWS**

[Enhancements and What's New in Hadoop 3.0](#) — A brief overview of incoming changes in Hadoop 3 and their benefits.

SHUBHAM SINHA **NEWS**

[How to migrate from RDBMS to MongoDB](#) — Learn why today's most innovative companies are migrating to MongoDB.

MONGODB **SPONSORED**

[The Many Faces of DISTINCT in Postgres](#)

HAKI BENITA

[Playing with SQL Server 2017's 'SQL Graph'](#)

DENIS GOBO

[10 Techniques to Boost Your Data Modeling](#)

SHELBY BLITZ

[The Motivation for Native Graph Databases](#)

JIM WEBBER

[Crunching Machine Learning And Databases Together on GPUs](#) — A look at accelerating relational DB processing with GPUs.

TIMOTHY PRICKETT MORGAN

[CIOs Growing Weary of Database Lock-In, Increasingly Buying Into Open Source](#) — Oracle is slipping in its hold on the database market, but it's going to take a long time to topple.

MATT ASAY

[JSONata: A Declarative Syntax for Querying JSON Data](#)

ANDREW COLEMAN **VIDEO**

[ToroDB: A Layer That Turns Postgres Into a NoSQL Database](#)

TORODB **TOOLS**

[SQLPad: Webapp to Run SQL Queries and Visualize Results](#) — Run SQL in your browser, chart the results. Supports MySQL, SQL Server and Vertica.

RICK BERGFALK **TOOLS**

[Manage all your Database-as-a-service needs in one place](#) — Manifold is your hub for amazing developer services, like JawsDB and RedisGreen

MANIFOLD **SPONSORED**

[SQL Notebook: A Windows App to Explore Tabular Data](#) — [Related Hacker News discussion.](#)

BRIAN LUFT **TOOLS**

[Trinity: A High Performance, C++ Search Library](#) — A building block for your own search engines.

MARK PAPADAKIS **CODE**

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).