Database Weekly

Issue 163 — July 14, 2017

BigchainDB 1.0: A Scalable Blockchain Database — BigchainDB is a decentralized database with blockchain characteristics, at scale. GitHub repo.
TIM DAUBENSCHÜTZ

Delivering Billions of Messages Exactly Once — The single requirement of all data pipelines is that they cannot lose data, but de-duplicating messages can be the trickier part.
SEGMENT

A pipeline for continuous data migration with the SQLizer API — With data holed up in third party apps everywhere, APIs are the sensible way to continuously migrate data to your own databases. The SQLizer API gives you essential conversion to SQL capability for your migration pipeline, all from your command line.
SQLIZER <mark>SPONSORED</mark>

Do We Need Specialized Graph Databases? — A look through an academic paper where the authors ran a graph data management system benchmark suite over several popular graph databases with interesting results.
THE MORNING PAPER

Working with DynamoDB At Scale is Very Tricky — Or, as the author puts it, *"You probably shouldn't use DynamoDB".* (Note: DynamoDB is AWS's NoSQL database service.)
JONO MACDOUGALL

ORMs Should Update "Changed" Values, Not Just "Modified" Ones — How the SQL language and its implementations distinguish between changed values and modified values.
LUKAS EDER

Using SQL Within Google Sheets — Did you know Google's Web-based spreadsheet tool includes a 'query' function for using SQL against your spreadsheet? This is a guide to using it.
BEN COLLINS

All About Indexes: The Very Basics — Well placed and defined indexes will improve your database's performance significantly. A popular article from last year.
EMIL DRKUŠIĆ

New CEO, Same Strategy for Couchbase
DATANAMI

Apache Spark 2.2.0 Released — The _structured streaming_ stream processing engine is now considered GA.
APACHE FOUNDATION

Estimating CPU Per Query with Weighted Linear Regression — This eBook explains how a specialization of linear regression answers seemingly unsolvable database questions.
VIVIDCORTEX `SPONSORED`

The Seven Types of SQL Joins
EREN BAYDEMIR

Modeling Polymorphic Associations in a Relational Database
JACK CHRISTENSEN

Access all of the content from MongoDB World 2017 — Watch the keynote videos and download all of the breakout session presentations here.
MONGO DB `VIDEO` `SPONSORED`

MongoDB Users Discuss Their NoSQL Journeys
DATANAMI

Enhanced Real-Time Data Processing with Redis Labs' In-Memory Database Solution
DATAVERSITY

BuntDB: Embeddable, In-Memory Key/Value Database — A low-level, in-memory (but persists to disk), key/value store in pure Go, with geospatial support. Inspired by BoltDB.
JOSH BAKER `CODE`