

[« Prev](#)[Next »](#)

## Database Weekly

Issue 181 — November 17, 2017

**SQL Operations Studio: A New Data Management Tool from Microsoft** — A Windows, macOS and Linux app for working with SQL Server, Azure SQL DB and SQL DW.

MICROSOFT **TOOLS**

**Elasticsearch 6.0.0 Released** — The popular full-text search-oriented database gains zero downtime upgrades, faster restarts, faster query times, and more. The new version is based on Lucene 7.

CLINTON GORMLEY

**Performance Analytics for Amazon Redshift. For Data Engineering Heroes** — Amazon Redshift is incredibly powerful, but you can spend a lot of time

fighting fires. We fix that. [intermix.io](#) gives you the tools you need to analyze your Redshift performance and improve the toolchain of everyone downstream from your cluster.

INTERMIX.IO **SPONSORED**

**Building a Microservices Ecosystem with Kafka Streams and KSQL** — How to build fine grained, streaming microservices with the Apache Kafka Streams API and KSQL.

BEN STOPFORD

**Datasette: Instantly Build an API for SQLite Databases** — Got an SQLite database you want to share on the web? This is the tool for you.

SIMON WILLISON **TOOLS**

**MariaDB Coming to Azure as Microsoft Joins the MariaDB Foundation** — The popular MySQL fork continues to make big steps.

ARS TECHNICA

**Distributed PostgreSQL Settling Into Cloud** — Citus brings horizontal scalability to Postgres users.

DATANAM

**Google Cloud's Distributed NoSQL Database Spanner Goes Global**

THE NEW STACK

**Azure Cosmos DB Launches Native Support for the Apache Cassandra API** —

Allows you to take Cassandra-backed apps to Azure using Microsoft's scalable multi-model database, Cosmos.

MICROSOFT AZURE

**Using Grafana with SiriDB** — Learn how to setup a Grafana dashboard using the SiriDB plugin.

TRANSEPTOR TECHNOLOGY **SPONSORED**

## **Troubleshooting MongoDB 100% CPU Load and Slow Queries**

IGOR KHOMENKO

**Faster Bulk Loading in Postgres with 'Copy'** — `\copy` is a mechanism for you to bulk load data in or out of Postgres.

CITUSDATA

**How ACID Transactions Work in a Globally Distributed Database** — .. or FaunaDB, specifically.

MATT FREELS

**The Mode Ordered-Set Aggregate Function** — A way to find the most frequent input value in a group.

DIMITRI FONTAINE

## **How to Avoid Excessive Sorts in Window Functions**

LUKAS EDER

## **How to Approach Cassandra NoSQL Data Model Design**

SHANNEN WEHRMAN

**Continuous Delivery: GoCD vs. Jenkins** — We compared GoCD with Jenkins on philosophy, continuous delivery & plugins. Understand how these CI/CD tools fit your needs.

GOCD **SPONSORED**

**PostgreSQL 10: An 'Amplified' Version of PostgreSQL** — Why version 10 is “*the defining version*” of the popular RDBMS.

COMMANDPROMPT.COM

**YCSB Benchmark Results for YugaByte and Apache Cassandra** — Yugabyte is a new database we featured 2 weeks ago which is compatible with Cassandra’s API.

YUGABYTE FORUM

## **How Should You Build A High-Performance Column Store for the 2020s?**

DANIEL LEMIRE

**dbmgr: CLI Tool to Back Up, Restore, and Provision Dev Databases** — For MySQL and Postgres only so far.

GITHUB.COM **TOOLS**

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).