

[« Prev](#)[Next »](#)

Database Weekly

Issue 187 — January 19, 2018

Datasette Publish: A Webapp for Publishing CSV Files as an Online Database — A tool that turns one or more CSV files into an online database with a JSON API, powered by SQLite and the [Zeit Now](#) deployment platform.

SIMON WILLISON [TOOLS](#)

Bistro: A Lightweight Column-Oriented Data Processing Engine — Written in Java, Bistro is an interesting new light-weight column-oriented data processing engine based on a novel column-oriented logical data model.

ALEXANDR SAVINOV [CODE](#)

Use SQL in MongoDB? But of Course You Can. We'll Show You How — And there's so much more to discover. Like SQL Import & Export. Or our

innovative Query Code that lets you instantly generate code from our Query Builder in Node.js, Python, or C#.

STUDIO 3T [SPONSORED](#)



SQL Keys in Depth — An attempt at boiling down the various arguments, pros and cons of different types of keys, whether they're auto-incrementing integers, UUIDs, or other things.

JOE NELSON

Fueled by Kafka, Stream Processing Poised for Growth — Once a niche technique used only by large organizations, stream processing has emerged as a popular way to deal with massive amounts of freshly generated data.

DATANAM

How Alibaba Created a Stream System with Millions of Feeds — A first hand look at how Chinese ecommerce company Alibaba scales using its own custom extensions to PostgreSQL.

ALIBABA TECH

Using Redis as a True Distributed, Fault-Tolerant Database — The 'trick' here is that you're not using Redis *per se*, but [Yugabyte](#), a database that *supports* the Redis protocol and can be used in place of it.

KARTHIK RANGANATHAN

Go Deep into Today's Modern Data Tier to See the Root of the Problem Fast

SELECTSTAR [SPONSORED](#)

How We Sped Up Eventually-Consistent Database Global Syncs with Caching —

They reduced an already solid time of 1 second for global consistency to 0.25 seconds by adding a volatile cache.

CHAITANYA PATCHAVA

Don't Throw Out Your Algorithms Book Just Yet: Classical Data Structures That Can Outperform Learned Indexes —

An interesting response to [this paper](#) we linked to last year about machine learning based indexing systems.

BAILIS, SHENG TAI, THAKER, AND ZAHARIA

Graphile: Instant GraphQL APIs for Postgres Databases — Auto-discovers tables, columns, relations, procedures and more, then runs a high-performance secure GraphQL API server that adheres to best practices.

BENJIE GILLAM [TOOLS](#)

What's The Effect of 'Meltdown' on Postgres's Performance? — How have the fixes for the recent Spectre and Meltdown CPU vulnerabilities affected Postgres performance? This first look suggests about a 7% hit.

2NDQUADRANT

How to Migrate an Azure SQL Database to Amazon Aurora

STEVE ABRAHAM

Using SymmetricDS to Replicate Your Database Structure — An open source database replication tool that can be used to keep heterogeneous databases in sync.

CHRIS HENSON

Analyzing Cassandra Performance with Flame Graphs

JON HADDAD

Linux Cloud Hosting Starting at 1GB of RAM for \$5/mo

LINODE CLOUD HOSTING [SPONSORED](#)

Scaling YugaByte DB to Millions of Reads and Writes — An open source database that supports Cassandra and Redis APIs.

BHARAT BADDEPUDI

PL/Swift: Write Custom PostgreSQL Functions in Swift

THE ALWAYS RIGHT INSTITUTE [TOOLS](#)

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).