

[« Prev](#)[Next »](#)

#220 — SEPTEMBER 7, 2018

[READ ON THE WEB](#)

Database Weekly

Select Star SQL: An Interactive SQL Tutorial — Want to get a bit better at SQL? This interactive online guide will give you some practical experience and is oriented, intriguingly, around querying a dataset of death row inmates..

SELECT STAR SQL

An Introduction to Druid, Interactive Analytics at Big Scale — Druid is an Apache-incubated high performance analytics data store for event-driven data (or, as the author describes it, *“a real-time columnar timeseries database on steroids”*).

FLORIAN ZYSSET

The Most Frequently Asked Questions: A Guide for the Overly Eager New InfluxDB User — This blog contains the most FAQ when you're getting started with InfluxDB, the fastest growing open source TSDB.

INFLUXDATA SPONSOR



'I Don't Want to Learn Your Garbage Query Language' — *“This is a bit of a rant but I really don't like software that invents its own query language.”* This opinion piece fuelled [an extensive discussion](#) on Hacker News.

ERIK BERNHARDSSON

Flink Delivers ACID Transactions on Streaming Data — The folks behind Apache Flink, an open source stream processing framework, just delivered something you may not have realized you needed: ACID guarantees in a stream processing framework.

DATANAMI

SQream DB 3.0 Boasts 15x Speedup for Its GPU Data Warehouse — [SQream](#) is a commercial, GPU-based data warehouse system aimed at bringing a relatively budget option to the multi-terabyte analytics market.

DATANAMI

Integrating MongoDB and Amazon Kinesis for Intelligent, Durable Streams — A tutorial looking at the MongoDB Stitch setup needed to both durably store data on MongoDB Atlas while simultaneously pushing it to Amazon Kinesis to perform stream-oriented analytics.

MICHAEL LYNN

ArangoDB 3.4 Release Candidate 1 Released: What's New? — ArangoDB is a highly available multi-model NoSQL database & 3.4 introduces [ArangoSearch](#), a search layer, and support for GeoJSON and Google's S2 Geo Index.

ARANGODB

Coming to Postgres 11: Parallel CREATE INDEX — Being able to create indexes in parallel is a significant development.

HANS-JÜRGEN SCHÖNIG

How to Download the Stack Overflow Database for SQL Server — The latest dump expands to a 304GB(!) SQL Server 2008 database. There's also a smaller 'laptop' build. An ideal dataset to stress or test SQL Server with.

BRENT OZAR

[Find the Next Non-NULL Row in a Series With SQL](#)

LUKAS EDER

[Migrating a Standalone PostgreSQL Instance to Amazon Aurora](#)

PAVEL TSIUKHTSIAYEU

[Break Free from Oracle](#) — Reduce database costs by 80% or more.

ENTERPRISEDB [SPONSOR](#)

[Spotify Open Sources 'cstar', Its Cassandra Orchestration Tool](#) — At 3000 nodes large, Spotify needed some tooling to help it manage its Cassandra cluster. This tool lets them (and now you) run arbitrary scripts on all hosts in a cluster in a “topology aware” fashion.

EMILIO DEL TESSANDORO (SPOTIFY)

[Sqorn: A Fast, Composable JavaScript SQL Query Builder](#) — Uses a functional style that models query compilation as a series of transformations. [Full tutorial here.](#)

SUFYAN DAWOODJEE

[Redis 5.0 RC5 Released](#)

SALVATORE SANFILIPPO

[« Prev](#)

[Next »](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).