

[« Prev](#)

#228 — NOVEMBER 2, 2018

[READ ON THE WEB](#)

Database Weekly



[Neo4j Gets the Largest Single Investment in the Graph Space](#) — The CEO of Neo4j, the company behind the eponymous graph database system, explains the \$80 million investment and its relationship to the future of graph technology and AI.

EMIL EIFREM (NEO4J)

[Gitbase: A Solution for Exploring Git Repos with SQL](#) — A look at an open source project that makes it possible to query code repositories with SQL.

FRANCESCA CAMPOY

[Kubernetes, Analytics, Pipeline as Code - What's New in GoCD](#) — The continuous delivery space is moving at the speed of light, and so is GoCD. Recently GoCD has released some exciting features including cloud native support, pipeline as code, CD analytics and more.

GOCD SPONSOR



[20+ MongoDB Alternatives You Should Know About](#) — These aren't exactly *drop in* replacements, but more a roundup of other databases (including PostgreSQL and TiDB) that can handle similar workloads if you're concerned about Mongo's recent license change.

PETER ZAITSEV (PERCONA)

[Rockset, an SQL Cloud Service, Emerges from Stealth](#) — An analytics startup launched by former Facebook engineers emerged this week with a cloud service that 'runs SQL directly on raw data'.

DATANAM

[Cockroach Labs Announces a Managed CockroachDB Service](#) — A new geo-distributed database as a service built around [CockroachDB](#), an open source SQL database system.

SPENCER KIMBALL

[TimescaleDB 1.0 is Now Production Ready](#) — We mentioned the first 1.0 release candidate [several weeks ago](#), but the time-series database built on PostgreSQL is now truly production ready.

MIKE FREEDMAN

[Querying Data Streams with Apache Pulsar SQL](#) — [Pulsar](#) is a distributed pub-sub system originally built at Yahoo. This article covers the motivation and use cases for the SQL querying of streaming data going through Pulsar.

JERRY PENG

[A Jepsen Report on MongoDB 3.6.4](#) — Jepsen performs analysis and research into distributed systems' safety, and here are the results of their testing of MongoDB's support for causal consistency and sharded clusters.

JEPSEN

[Visualizing SQL Plan Execution Time with FlameGraphs](#) — This example uses Oracle and Postgres for example purposes but the idea can be adapted to any database that can output the right stats.

TANEL PODER

[Pivoting with SQL in Apache Spark 2.4](#) — A look at how a common operation to 'rotate' tables by turning unique values from one column into individual columns can be achieved in Spark using SQL. [A more generic guide here.](#)

MARYANN XUE

[Recreating a Column Oriented Database in PostgreSQL](#) — This is *rather niche*, but you may find this experiment in making Postgres act more like a column oriented DB interesting.

BRIAN DAMS

[How Streak Built a Graph Database on Google Cloud Spanner to Wrangle Billions of Emails](#)

GOOGLE CLOUD BLOG

[Webinar: MongoDB Mobile and MongoDB Stitch – Introduction and Latest Developments](#)

MONGODB SPONSOR

[SQL Trainer: Learn SQL by Doing Live Data Exercises](#)

BYTESCOUT

[Analytics Zoo: Distributed Tensorflow, Keras and BigDL on Apache Spark](#) — Provides a unified analytics and AI platform that unites Spark, TensorFlow, Keras and BigDL into an integrated pipeline.

INTEL CORPORATION

□ Jobs

[UK Tech Jobs? Check Out hackajob](#) — Get matched with companies like Apple, Skyscanner and KPMG based on your skills. They apply to you.

HACKAJOB

[« Prev](#)

Want to subscribe? Enter your address here

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's [our privacy policy](#).